

# Credit Card Fraud Detection

**Dataset Link :** <https://www.kaggle.com/datasets/mlg-ulb/creditcardfraud>

## Dataset Description:

The dataset contains transactions made by credit cards over a two-day period in **September 2013** by European cardholders. It comprises a total of **284,807** transactions, out of which only **492** transactions are labeled as fraudulent. The dataset is highly imbalanced, with the majority of transactions being non-fraudulent.

The dataset includes the following attributes:

**Time:** The number of seconds elapsed between this transaction and the first transaction in the dataset.

**V1 to V28:** Features resulting from a PCA transformation for confidentiality reasons.

**Amount:** The transaction amount.

**Class:** The target variable indicating whether the transaction is fraudulent (1) or not (0).

Glmpse of the Dataset ->

## Import dataset(.csv file) into Rstudio:

```
card=read.csv("C:/Users/O M A R/Downloads/Document/Data Science/PomPom/Book1.csv")
```

```
print(card)
```

```
> card=read.csv("C:/Users/O M A R/Downloads/Document/Data Science/PomPom/Book1.csv")
> print(card)
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11
1	0	-1.3598071	-0.07278117	2.53634674	1.37815522	-0.338320770	0.46238778	0.239598554	0.098697901	0.3637870	0.09079417	-0.55159953
2	0	1.1918571	0.26615071	0.16648011	0.44815408	0.060017649	-0.08236081	-0.078802983	0.085101655	-0.2554251	-0.16697441	1.61272666
3	1	-1.3583541	-1.34016307	1.77320934	0.37977959	-0.503198133	1.80049938	0.791460956	0.247675787	-1.5146543	0.20764287	0.62450146
4	1	-0.9662717	-0.18522601	1.79299334	-0.86329128	-0.010308880	1.24720317	0.237608940	0.377435875	-1.3870241	-0.05495192	-0.22648726
5	2	-1.1582331	0.87773676	1.54871785	0.40303393	-0.407193377	0.09592146	0.592940745	-0.270532677	0.8177393	0.75307443	-0.82284288

This code segment loads the credit card fraud dataset using the read.csv function and displays the first few rows of the dataset to understand its structure.

### Data Processing:

```
card$Amount <- scale(card$Amount)
card$Time <- scale(card$Time)
set.seed(42)
index <- sample(1:nrow(card), 0.7 * nrow(card))
train_data <- card[index, ]
test_data <- card[-index, ]
```

### Overall summary of the dataset:

summary(card)

```
> summary(card)
```

	Time	V1	V2	V3	V4	V5	V6	V7
Min.	: 0.0	Min. :-6.0932	Min. :-12.1142	Min. :-5.6950	Min. :-4.516	Min. :-6.6320	Min. :-2.1457	Min. :-2.7054
1st Qu.:	93.0	1st Qu. :-0.8976	1st Qu. :-0.1687	1st Qu. : 0.3121	1st Qu. : 1.179	1st Qu. :-1.3602	1st Qu. : 0.3757	1st Qu. :-1.1348
Median :	215.0	Median :-0.3655	Median : 0.2858	Median : 0.8922	Median : 1.810	Median :-1.3602	Median : 0.9539	Median :-1.1348
Mean :	215.4	Mean :-0.1702	Mean : 0.2114	Mean : 0.8729	Mean : 1.312	Mean :-0.8997	Mean : 0.7448	Mean :-0.6603
3rd Qu.:	326.0	3rd Qu. : 1.1108	3rd Qu. : 0.8766	3rd Qu. : 1.5138	3rd Qu. : 1.810	3rd Qu. :-0.5143	3rd Qu. : 0.9539	3rd Qu. :-0.1074
Max.	:450.0	Max. : 1.5861	Max. : 5.2674	Max. : 3.7729	Max. : 4.076	Max. : 3.2820	Max. : 5.1221	Max. : 4.8084

### Number of Column and Row:

```
sprintf("Rows: %d Columns: %d",nrow(card), length(names(card)))
```

```
> sprintf("Rows: %d Columns: %d",nrow(card), length(names(card)))
[1] "Rows: 599 Columns: 31"
```

### Count of the Missing value from each attribute:

```
> colSums(is.na(card))
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14	V15	V16	V17	V18	V19
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
V20	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class										
0	0	0	0	0	0	0	0	0	0	0										

### Conversion of all Non-Numeric columns to Numeric (excluding 'Class'):

```
> numeric_cols <- sapply(card, is.numeric)
> card[, numeric_cols] <- lapply(card[, numeric_cols], as.numeric)
```

### Handle missing values (replace NAs with zeros in numeric columns):

```
> numeric_cols <- sapply(card, is.numeric)
> card[, numeric_cols] <- lapply(card[, numeric_cols], function(col) {
+   col[is.na(col)] <- 0
+   return(col)
+ })
```

## Data Normalize the features (except 'Class' and 'Time'):

```
> cols_to_normalize <- colnames(card)[!(colnames(card) %in% c("class", "Time"))]
> card[, cols_to_normalize] <- scale(card[, cols_to_normalize])
```


## Crosscheck of the Changes:

```
> head(card)
  Time      V1      V2      V3      V4      V5      V6      V7      V8      V9      V10      V11      V12
1    0 -0.8958854 -0.23701777 1.6238181 0.06285398 0.5758058 -0.3286463 1.0942915 -0.49272197 -0.8768982 0.5653416 0.4129761 -2.35926538
2    0  1.0258322  0.04570169 -0.6895607 -0.81905543 0.9844055 -0.9626256 0.7071017 -0.52016512 -1.4842311 0.1255089 2.7186744 2.65307108
3    1 -0.8947911 -1.29420227 0.8788708 -0.88389417 0.4066812 1.2286500 1.7653797 -0.19201968 -2.7193027 0.7647216 1.6658991 -0.32255370
4    1 -0.5995048 -0.33081344 0.8981832 -2.06268400 0.9122674 0.5847230 1.0918721 0.06989272 -2.5941210 0.3166538 0.7593244 0.01142972
5    2 -0.7440753 0.55585517 0.6597302 -0.86184234 0.5051591 -0.7551404 1.5239707 -1.23799017 -0.4316546 1.6953961 0.1240153 1.08346690
6    2 -0.1925879 0.62491121 0.2618374 -1.40358631 1.3546733 -0.9013709 1.3820101 -0.16650959 -1.7914679 -0.2233165 2.4294779 0.55245741

  V13      V14      V15      V16      V17      V18      V19      V20      V21      V22      V23      V24
1 0.1110417 0.4852478 2.48529475 0.1822639 -0.4419577 0.89939729 1.079741298 0.8477036 -0.01856707 0.2613796 -0.1296688 -0.02326148
2 1.7392184 0.8100350 1.52054285 1.8883624 -1.0714823 0.4751484 -0.005730137 0.1141520 -0.77944234 -2.0588876 0.6794533 -1.15111607
3 1.9901809 0.7670133 3.50226802 -4.2361663 1.3172504 0.5993067 -4.183682497 1.4738469 0.95808656 1.5116067 3.7672188 -2.11998514
4 1.7597420 0.5303494 0.05250299 -0.8937225 -2.1817909 4.7787725 -2.151572462 -0.2038883 -0.34861221 -0.4286528 -0.4347551 -3.46832255
5 2.6814443 -1.0834210 0.98703653 0.2168697 -1.3098702 0.7658429 1.868497604 1.2073443 0.01398530 1.5789459 -0.232727 0.18285633
6 0.8075182 0.8229153 1.38388479 1.7747990 -0.9609526 0.9798048 0.216564794 0.4667445 -0.71518268 -1.8592754 0.1915785 -1.23867752

  V25      V26      V27      V28      Amount      class
1 -0.6917473 -0.19347601 0.25865658 -0.202363541 0.43034713 0
2 -0.5522222 0.80544967 -0.36548981 -0.004615734 -0.36073302 0
3 -2.3393523 -0.03486316 -0.56852880 -0.416259422 1.66351256 0
4 1.1821526 -0.29753243 -0.05151096 0.253689374 0.28971545 0
5 -1.9000490 1.99904374 0.63462865 1.103195005 0.00161432 0
6 -1.9967865 0.74209191 0.78535225 0.362147533 -0.35545664 0
```

## Correlations : Pearson Correlation



**Confusion Matrix and Accuracy calculated:**



```
confusion_matrix <- confusionMatrix(knn_model, test_data$Class)

print(confusion_matrix)
```

```
> confusion_matrix <- confusionMatrix(knn_model, test_data$Class)
> print(confusion_matrix)
Confusion Matrix and Statistics

          Reference
Prediction 0    1
0  167    5
1    2    6

      Accuracy : 0.9611
      95% CI   : (0.9215, 0.9842)
    No Information Rate : 0.9389
    P-Value [Acc > NIR] : 0.1351

      Kappa : 0.6116

  Mcnemar's Test P-Value : 0.4497

      Sensitivity : 0.9882
      Specificity : 0.5455
    Pos Pred Value : 0.9709
    Neg Pred Value : 0.7500
      Prevalence : 0.9389
    Detection Rate : 0.9278
    Detection Prevalence : 0.9556
    Balanced Accuracy : 0.7668

    'Positive' Class : 0
```

### 10-fold cross validation:

**To perform 10-fold cross-validation on dataset using the K-Nearest Neighbors (KNN) classifier, I can use the following code:**

```
card$Class <- as.factor(card$Class)

num_folds <- 10

set.seed(123)

cv <- createFolds(card$Class, k = num_folds)

ctrl <- trainControl(method = "cv", number = num_folds)

knn_model <- train(Class ~ ., data = card, method = "knn",
  trControl = ctrl, tuneLength = 1,
  preProcess = c("center", "scale"),
  metric = "Accuracy")

cat("Mean Accuracy:", knn_model$results$Accuracy, "\n")
```

This code performs 10-fold cross-validation using the KNN classifier. It creates a cross-validation object using the `createFolds()` function, iterates through each fold, fits the KNN model, calculates accuracy for each fold, and finally calculates the mean accuracy across all folds.

```
> cat("Mean Accuracy:", knn_model$results$Accuracy, "\n")  
Mean Accuracy: 0.9332685
```

### Modeling Evaluation:

Applying machine learning techniques, such as K-nearest neighbors (KNN), helps classify transactions as fraudulent or non-fraudulent. To assess model performance, we utilized accuracy, recall, and precision metrics. The challenge lies in optimizing the trade-off between recall and precision due to the imbalanced dataset.

### Discussion:

The Credit Card Fraud Detection dataset provides a rich context for exploring the challenges and intricacies of detecting fraudulent transactions in financial systems. The dataset's complexity stems from its imbalanced nature, anonymized features, and the critical implications of accurate fraud detection. Let's delve into the key discussion points surrounding this dataset:

- **Key Insights:** Imbalanced nature, Feature anonymization, Importance of 'Time' and 'Amount'
- **Data Preparation:** Conversion to Numeric, Handling missing value, Feature normalization
- **Modeling and Evaluation:** K-nearest neighbors (KNN), 10-fold cross validation, Confusion Matrix and predict Accuracy

### Conclusion:

**The Credit Card Fraud Detection** dataset poses a real-world problem with significant implications for financial security. The imbalance, anonymized features, and critical nature of the task require careful preprocessing, thoughtful feature engineering, and advanced modeling techniques. Achieving high accuracy on this dataset may be misleading due to the dominant non-fraudulent class. As such, a comprehensive evaluation should involve considering recall and precision, striking a balance between detecting fraud and minimizing false positives. In practice, deploying and maintaining such models require ongoing monitoring and adjustment to adapt to evolving fraudulent behaviors.